

Decision Vitals

Vital signs for the decisions you've already made. Decision Vitals watches the assumptions a business decision rests on and tells you when the evidence starts to turn against one.

The problem

Most decisions get made once and are rarely revisited. The assumptions underneath them, about customers, capacity, timing, and the market, quietly go out of date, and nobody notices until the damage shows up in the numbers. Most teams monitor their decisions by gut feel: the review happens after the failure, not before it.

The idea


Every decision rests on a few assumptions. Some are **critical**: if they turn out to be wrong, the decision could seriously weaken or break. Others are **supporting**: they still matter, but the decision can probably survive if they change. For each one there is a **warning signal**, the specific thing you should watch for that would show the assumption may no longer be true. This way of pressure-testing a plan comes from a decades-old risk-planning method (credited below); the bet here is that AI agents make it light enough to actually use on everyday decisions.

What it does

You record a decision and Decision Vitals identifies the assumptions underneath it, marks which are critical, and gives each a warning signal to watch. You log evidence as it accumulates: meeting notes, tickets, customer feedback, market updates. When you review the decision, specialized AI agents examine the evidence from different perspectives and combine their findings into a Decision Health Report: an overall grade, an assessment of each assumption with its supporting evidence, the strongest case against it, and recommended next steps. Reviews are numbered and dated, so a decision builds up a health history instead of a single one-time assessment.

How reviews work

Each review is carried out by six specialized AI agents, each with a single job: two identify and classify the assumptions when you record a decision, and four run the review. Different agents examine the evidence, make the case against each assumption, assess where each one stands, and write the report. Splitting the work this way keeps every step focused, and it means no single perspective dominates: the agent arguing against an assumption has no say in the final grade.



Every step's findings stay attached to the result, so each conclusion can be traced back to the evidence and reasoning behind it. A few fixed rules keep the results consistent: a critical assumption with strong evidence against it can never be marked as still holding, and the overall health grade is computed from the individual assessments.

Decision Vitals also connects to Claude directly: from a conversation, Claude can list your decisions, inspect their assumptions, and log new evidence, which shows up in the app ready for the next review.

Product decisions

Paste-only evidence	No integrations, no file parsing. Adoption beats automation for a tool like this: if logging a piece of evidence takes more than a paste, it doesn't happen.
You edit the assumptions	Before the first review you can reword, rerank, or delete any assumption, so human judgment goes in before the machine's, not after.
Numbered reviews, not silent monitoring	Each review is an explicit, dated event with an evidence record that doesn't change after the fact. That's the honest, lightweight version of "always watching."
Demo mode	The public site replays real recorded reviews, labeled as such, so anyone can try it at no cost; live reviews are passphrase-protected.

Limitations

Results are only as good as the evidence you log, and sources aren't weighted, a rumor and an audited number count the same. It works on one decision at a time, with no links between related decisions. And the model's read on how strongly a piece of evidence cuts is a judgment call, the built-in rules limit the worst mistakes but don't remove them.

What's next

Pulling evidence automatically from one source (a project tool is the natural first one), linking related decisions so a broken assumption in one flags the others, and alerting the owner the moment a warning signal shows up rather than waiting for a manual review.

Built with Claude Code on the web; agents hosted on Claude Managed Agents; deployed on Vercel. Inspired by [Assumption-Based Planning \(RAND\)](#).